

# Speed: The GCS ENCS Cluster

Serguei A. Mokhov      Gillian A. Roper      Carlos Alarcón Meza  
Network, Security and HPC Group\*  
Gina Cody School of Engineering and Computer Science  
Concordia University  
Montreal, Quebec, Canada  
`rt-ex-hpc@encs.concordia.ca`  
**Version 7.2-dev-02**

## Abstract

This document presents a quick start guide to the usage of the Gina Cody School of Engineering and Computer Science compute server farm called “Speed” – the GCS Speed cluster, managed by the HPC/NAG group of the Academic Information Technology Services (AITS) at GCS, Concordia University, Montreal, Canada.

## Contents

<b>1</b>	<b>Introduction</b>	<b>3</b>
1.1	Resources . . . . .	3
1.2	Team . . . . .	3
1.3	What Speed Consists of . . . . .	3
1.4	What Speed Is Ideal For . . . . .	4
1.5	What Speed Is Not . . . . .	5
1.6	Available Software . . . . .	5
1.7	Requesting Access . . . . .	6
<b>2</b>	<b>Job Management</b>	<b>7</b>
2.1	Getting Started . . . . .	7
2.1.1	SSH Connections . . . . .	7
2.1.2	Environment Set Up . . . . .	7
2.2	Job Submission Basics . . . . .	8
2.2.1	Directives . . . . .	8
2.2.2	Module Loads . . . . .	10
2.2.3	User Scripting . . . . .	10
2.3	Sample Job Script . . . . .	11
2.4	Common Job Management Commands Summary . . . . .	12
2.5	Advanced <code>sbatch</code> Options . . . . .	14
2.6	Array Jobs . . . . .	14
2.7	Requesting Multiple Cores (i.e., Multithreading Jobs) . . . . .	15
2.8	Interactive Jobs . . . . .	16
2.8.1	Command Line . . . . .	16
2.8.2	Graphical Applications . . . . .	16
2.8.3	Jupyter Notebooks in Singularity . . . . .	17
2.8.4	Jupyter Labs in Conda and Pytorch . . . . .	18
2.8.5	VScode . . . . .	21

---

\*The group acknowledges the initial manual version VI produced by Dr. Scott Bunnell while with us as well as Dr. Tariq Daradkeh for his instructional support of the users and contribution of examples.

2.9	Scheduler Environment Variables . . . . .	21
2.10	SSH Keys For MPI . . . . .	22
2.11	Creating Virtual Environments . . . . .	23
2.11.1	Anaconda . . . . .	23
2.11.2	Python . . . . .	24
2.12	Example Job Script: Fluent . . . . .	25
2.13	Example Job: efficientdet . . . . .	25
2.14	Java Jobs . . . . .	26
2.15	Scheduling On The GPU Nodes . . . . .	26
2.15.1	P6 on Multi-GPU, Multi-Node . . . . .	27
2.15.2	CUDA . . . . .	28
2.15.3	Special Notes for sending CUDA jobs to the GPU Queue . . . . .	28
2.15.4	OpenISS Examples . . . . .	28
2.16	Singularity Containers . . . . .	29
<b>3</b>	<b>Conclusion</b>	<b>30</b>
3.1	Important Limitations . . . . .	30
3.2	Tips/Tricks . . . . .	31
3.3	Use Cases . . . . .	31
<b>A</b>	<b>History</b>	<b>32</b>
A.1	Acknowledgments . . . . .	32
A.2	Migration from UGE to SLURM . . . . .	33
A.3	Phases . . . . .	34
A.3.1	Phase 4 . . . . .	34
A.3.2	Phase 3 . . . . .	34
A.3.3	Phase 2 . . . . .	35
A.3.4	Phase 1 . . . . .	35
<b>B</b>	<b>Frequently Asked Questions</b>	<b>35</b>
B.1	Where do I learn about Linux? . . . . .	35
B.2	How to use the “bash shell” on Speed? . . . . .	35
B.2.1	How do I set bash as my login shell? . . . . .	35
B.2.2	How do I move into a bash shell on Speed? . . . . .	36
B.2.3	How do I use the bash shell in an interactive session on Speed? . . . . .	36
B.2.4	How do I run scripts written in bash on Speed? . . . . .	36
B.3	How to resolve “Disk quota exceeded” errors? . . . . .	36
B.3.1	Probable Cause . . . . .	36
B.3.2	Possible Solutions . . . . .	36
B.3.3	Example of setting working directories for COMSOL . . . . .	37
B.3.4	Example of setting working directories for Python Modules . . . . .	37
B.4	How do I check my job’s status? . . . . .	38
B.5	Why is my job pending when nodes are empty? . . . . .	38
B.5.1	Disabled nodes . . . . .	38
B.5.2	Error in job submit request. . . . .	39
<b>C</b>	<b>Sister Facilities</b>	<b>39</b>
	<b>Annotated Bibliography</b>	<b>40</b>

## 1 Introduction

This document contains basic information required to use “Speed” as well as tips and tricks, examples, and references to projects and papers that have used Speed. User contributions of sample jobs and/ or references are welcome. Details are sent to the `hpc-m1` mailing list.

**Note:** On October 20, 2023 with workshops prior, we have completed migration to SLURM (see Figure 2) from Grid Engine (UGE/AGE) as our job scheduler, so this manual has been ported to use SLURM’s syntax and commands. If you are a long-time GE user, see Appendix A.2 key highlights of the move needed to translate your GE jobs to SLURM as well as environment changes. These are also elaborated throughout this document and our examples as well in case you desire to re-read it.

If you wish to cite this work in your acknowledgements, you can use our general DOI found on our GitHub page <https://dx.doi.org/10.5281/zenodo.5683642> or a specific version of the manual and scripts from that link individually.

### 1.1 Resources

- Our public GitHub page where the manual and sample job scripts are maintained (pull-requests (PRs), subject to review, are welcome):  
<https://github.com/NAG-DevOps/speed-hpc>  
<https://github.com/NAG-DevOps/speed-hpc/pulls>
- PDF version of this manual:  
<https://github.com/NAG-DevOps/speed-hpc/blob/master/doc/speed-manual.pdf>  
HTML version of this manual:  
<https://nag-devops.github.io/speed-hpc/>
- Our official Concordia page for the “Speed” cluster:  
<https://www.concordia.ca/ginacody/aits/speed.html>  
which includes access request instructions.
- All Speed users are subscribed to the `hpc-m1` mailing list.

### 1.2 Team

Speed is supported by:

- Serguei Mokhov, PhD, Manager, Networks, Security and HPC, AITS
- Gillian Roper, Senior Systems Administrator, HPC, AITS
- Carlos Alarcón Meza, Systems Administrator, HPC and Networking, AITS

We receive support from the rest of AITS teams, such as NAG, SAG, FIS, and DOG.  
<https://www.concordia.ca/ginacody/aits.html>

### 1.3 What Speed Consists of

- Twenty four (24) 32-core compute nodes, each with 512 GB of memory and approximately 1 TB of local volatile-scratch disk space (pictured in Figure 1).

- Twelve (12) NVIDIA Tesla P6 GPUs, with 16 GB of memory (compatible with the CUDA, OpenGL, OpenCL, and Vulkan APIs).
- 4 VIDPRO nodes, with 6 P6 cards, and 6 V100 cards (32GB), and 256GB of RAM.
- 7 new SPEED2 servers with 64 CPU cores each 4x A100 80 GB GPUs, partitioned into 4x 20GB each; larger local storage for TMPDIR.
- One AMD FirePro S7150 GPU, with 8 GB of memory (compatible with the Direct X, OpenGL, OpenCL, and Vulkan APIs).

## Current State: Physical




Figure 1: Speed

### 1.4 What Speed Is Ideal For

- To design and develop, test and run parallel, batch, and other algorithms, scripts with partial data sets. “Speed” has been optimised for compute jobs that are multi-core aware, require a large memory space, or are iteration intensive.
- Prepare them for big clusters:
  - Digital Research Alliance of Canada (Calcul Quebec and Compute Canada)
  - Cloud platforms
- Jobs that are too demanding for a desktop.
- Single-core batch jobs; multithreaded jobs typically up to 32 cores (i.e., a single machine).
- Multi-node multi-core jobs (MPI).




Figure 2: Speed SLURM Architecture

- Anything that can fit into a 500-GB memory space and a **scratch** space of approximately 10 TB.
- CPU-based jobs.
- CUDA GPU jobs (`speed-01|-03|-05`, `speed-17`, `speed-37-speed-43`).
- Non-CUDA GPU jobs using OpenCL (`speed-19` and `-01|03|05|17|25|27|37-43`).

## 1.5 What Speed Is Not

- Speed is not a web host and does not host websites.
- Speed is not meant for Continuous Integration (CI) automation deployments for Ansible or similar tools.
- Does not run Kubernetes or other container orchestration software.
- Does not run Docker. (**Note:** Speed does run Singularity and many Docker containers can be converted to Singularity containers with a single command. See Section 2.16.)
- Speed is not for jobs executed outside of the scheduler. (Jobs running outside of the scheduler will be killed and all data lost.)

## 1.6 Available Software

We have a great number of open-source software available and installed on “Speed” – various Python, CUDA versions, C++/Java compilers, OpenGL, OpenFOAM, OpenCV, TensorFlow,

OpenMPI, OpenISS, MARF [24], etc. There are also a number of commercial packages, subject to licensing contributions, available, such as MATLAB [13, 23], Abaqus [1], Ansys, Fluent [2], etc.

To see the packages available, run `ls -al /encs/pkg/` on `speed.encs`. In particular, there are over 2200 programs available in `/encs/bin` and `/encs/pkg` under Scientific Linux 7 (EL7). We are building an equivalent array of programs for the EL9 SPEED2 nodes.

- Popular concrete examples:
  - MATLAB (R2016b, R2018a, R2018b, ...)
  - Fluent (19.2, ...)
  - Singularity containers (see Section 2.16) can run other operating systems and Linux distributions, like Ubuntu’s, as well as converted Docker containers.
- We do our best to accommodate custom software requests. Python environments can use user-custom installs from within the scratch directory.
- A number of specific environments are available and can be loaded using the `module` command:
  - Python (2.3.x - 3.11.x)
  - Gurobi (7.0.1, 7.5.0, 8.0.0, 8.1.0)
  - Ansys (16, 17, 18, 19)
  - OpenFOAM (2.3.1, 3.0.1, 5.0, 6.0)
  - Cplex 12.6.x to 12.8.x
  - OpenMPI 1.6.x, 1.8.x, 3.1.3

## 1.7 Requesting Access

After reviewing the “What Speed is” (Section 1.4) and “What Speed is Not” (Section 1.5), request access to the “Speed” cluster by emailing: `rt-ex-hpc AT encs.concordia.ca`. GCS ENCS faculty and staff may request access directly. Students must include the following in their message:

- GCS ENCS username
- Name and email (CC) of the supervisor or instructor
- Written request from the supervisor or instructor for the ENCS username to be granted access to “Speed”

Non-GCS faculty / students need to get a “sponsor” within GCS, such that your guest GCS ENCS account is created first. A sponsor can be any GCS Faculty member you collaborate with. Failing that, request the approval from our Dean’s Office; via our Associate Deans Drs. Eddie Hoi Ng or Emad Shihab. External entities to Concordia who collaborate with GCS Concordia researchers, should also go through the Dean’s office for approvals. Non-GCS students taking a GCS course do have their GCS ENCS account created automatically, but still need the course instructor’s approval to use the service.

## 2 Job Management

In these instructions, anything bracketed like so, <>, indicates a label/value to be replaced (the entire bracketed term needs replacement). We use SLURM as the Workload Manager. It supports primarily two types of jobs: batch and interactive. Batch jobs are used to run unattended tasks.

TL;DR: Job instructions in a script start with **#SBATCH** prefix, for example:

```
#SBATCH --account=speed1 --mem=100M -t 600 -J job-name
#SBATCH --gpus=2 --mail-type=ALL -t 600 --mail-user=YOUR_USERNAME
```

We use **srun** for every complex compute step inside the script. Use interactive jobs to set up virtual environments, compilation, and debugging. **salloc** is preferred; allows multiple steps. **srun** can start interactive jobs as well (see Section 2.8). Required and common job parameters: job-name (J), mail-type, mem, ntasks (n), cpus-per-task, account, -p (partition).

### 2.1 Getting Started

Before getting started, please review the “What Speed is” (Section 1.4) and “What Speed is Not” (Section 1.5). Once your GCS ENCS account has been granted access to “Speed”, use your GCS ENCS account credentials to create an SSH connection to **speed** (an alias for **speed-submit.encs.concordia.ca**). All users are expected to have a basic understanding of Linux and its commonly used commands (see Appendix B.1 for resources).

#### 2.1.1 SSH Connections

Requirements to create connections to Speed:

1. An active **GCS ENCS user account**, which has permission to connect to Speed (see Section 1.7).
2. If you are off campus, an active connection to Concordia’s VPN. Accessing Concordia’s VPN requires a Concordia **netname**.
3. Windows systems require a terminal emulator such as PuTTY, Cygwin, or MobaXterm.
4. macOS systems do have a Terminal app for this or **xterm** that comes with XQuartz.

Open up a terminal window and type in the following SSH command being sure to replace <ENCSusername> with your ENCS account’s username.

```
ssh <ENCSusername>@speed.encs.concordia.ca
```

Read the AITS FAQ: How do I securely connect to a GCS server?

#### 2.1.2 Environment Set Up

After creating an SSH connection to Speed, you will need to make sure the **srun**, **sbatch**, and **salloc** commands are available to you. Type the command name at the command prompt and press enter. If the command is not available, e.g., (“command not found”) is returned, you need to make sure your **\$PATH** has **/local/bin** in it. To view your **\$PATH** type **echo \$PATH** at the prompt.

The next step is to copy a job template to your home directory and to set up your cluster-specific storage. Execute the following command from within your home directory. (To move to your home directory, type `cd` at the Linux prompt and press **Enter**.)

```
cp /home/n/nul-uge/template.sh . && mkdir /speed-scratch/$USER
```

**Tip:** the default shell for GCS ENCS users is `tcsh`. If you would like to use `bash`, please contact `rt-ex-hpc AT encs.concordia.ca`.

**Note:** If a “command not found” error appears after you log in to speed, your user account many have probably have defunct Grid Engine environment commands. See Appendix A.2 to learn how to prevent this error on login.

## 2.2 Job Submission Basics

Preparing your job for submission is fairly straightforward. Start by basing your job script on one of the examples available in the `src/` directory of our GitHub’s (<https://github.com/NAG-DevOps/speed-hpc>). Job scripts are broken into four main sections:

- Directives
- Module Loads
- User Scripting

You can clone the tip of our repository to get the examples to start with or download them individually via a browser or command line:

```
git clone --depth=1 https://github.com/NAG-DevOps/speed-hpc.git
cd speed-hpc/src
```

Then to quickly run some sample jobs, you can:

```
sbatch -p ps -t 10 bash.sh
sbatch -p ps -t 10 env.sh
sbatch -p ps -t 10 manual.sh
sbatch -p pg -t 10 lambdal-singularity.sh
```

### 2.2.1 Directives

Directives are comments included at the beginning of a job script that set the shell and the options for the job scheduler. The shebang directive is always the first line of a script. In your job script, this directive sets which shell your script’s commands will run in. On “Speed”, we recommend that your script use a shell from the `/encs/bin` directory.

To use the `tcsh` shell, start your script with `#!/encs/bin/tcsh`. For `bash`, start with `#!/encs/bin/bash`. Directives that start with `#SBATCH`, set the options for the cluster’s Slurm job scheduler. The script template, `template.sh`, provides the essentials:

```
#SBATCH --job-name=<jobname>          ## or -J. Give the job a name
#SBATCH --mail-type=<type>            ## Set type of email notifications
#SBATCH --mail-user=<YOUR_USERNAME>@encs.concordia.ca
#SBATCH --chdir=<directory>           ## or -D, Set working directory where output files will go
#SBATCH --nodes=1                     ## or -N, Node count required for the job
```



```
#SBATCH --ntasks=1                ## or -n, Number of tasks to be launched
#SBATCH --cpus-per-task=<corecount> ## or -c, Core count requested, e.g. 8 cores
#SBATCH --mem=<memory>            ## Assign memory for this job, e.g., 32G memory per node
```

Replace the following to adjust the job script for your project(s)

1. `<jobname>` with a job name for the job
2. `<YOUR_USERNAME>` with your GCS username
3. `<directory>` with the fullpath to your job's working directory, e.g., where your code, source files and where the standard output files will be written to. By default, `--chdir` sets the current directory as the job's working directory
4. `<type>` with the type of e-mail notifications you wish to receive. Valid options are: NONE, BEGIN, END, FAIL, REQUEUE, ALL
5. `<corecount>` with the degree of multithreaded parallelism (i.e., cores) allocated to your job. Up to 32 by default.
6. `<memory>` with the amount of memory, in GB, that you want to be allocated per node. Up to 500 depending on the node. NOTE: All jobs MUST set a value for the `--mem` option.

Example with short option equivalents:

```
#SBATCH -J tmpdir                ## Job's name set to 'tmpdir'
#SBATCH --mail-type=ALL          ## Receive all email type notifications
#SBATCH --mail-user=a_user@encs.concordia.ca
#SBATCH -D ./                    ## Use current directory as working directory
#SBATCH -N 1                     ## Node count required for the job
#SBATCH -n 1                     ## Number of tasks to be launched
#SBATCH -c 1                     ## Request 8 cores
#SBATCH --mem=32G                ## Allocate 32G memory per node
```

If you are unsure about memory footprints, err on assigning a generous memory space to your job, so that it does not get prematurely terminated. You can refine `--mem` values for future jobs by monitoring the size of a job's active memory space on `speed-submit` with:

```
sacct -j <jobID>
sstat -j <jobID>
```

This can be customized to show specific columns:

```
sacct -o jobid,maxvmsize,ntasks%7,tresusageouttot%25 -j <jobID>
sstat -o jobid,maxvmsize,ntasks%7,tresusageouttot%25 -j <jobID>
```

Memory-footprint values are also provided for completed jobs in the final e-mail notification as "maxvmsize". *Jobs that request a low-memory footprint are more likely to load on a busy cluster.*

Other essential options are `--time`, or `-t`, and `--account`, or `-A`.

- **--time=<time>** – is the estimate of wall clock time required for your job to run. As previously mentioned, the maximum is 7 days for batch and 24 hours for interactive jobs. Jobs with a smaller **time** value will have a higher priority and may result in your job being scheduled sooner.
- **--account=<name>** – specifies which Account, aka project or association, that the Speed resources your job uses should be attributed to. When moving from GE to SLURM users most users were assigned to Speed’s two default accounts **speed1** and **speed2**. However, users that belong to a particular research group or project are will have a default Account like the following **aits**, **vidpro**, **gipsy**, **ai2**, **mpackir**, **cmos**, among others.

### 2.2.2 Module Loads

As your job will run on a compute or GPU “Speed” node, and not the submit node, any software that is needed must be loaded by the job script. Software is loaded within the script just as it would be from the command line.

To see a list of which modules are available, execute the following from the command line on **speed-submit**.

```
module avail
```

To list for a particular program (**matlab**, for example):

```
module -t avail matlab
```

Which, of course, can be shortened to match all that start with a particular letter:

```
module -t avail m
```

Insert the following in your script to load the **matlab/R2020a**) module:

```
module load matlab/R2020a/default
```

Use, **unload**, in place of, **load**, to remove a module from active use.

To list loaded modules:

```
module list
```

To purge all software in your working environment:

```
module purge
```

Typically, only the **module load** command will be used in your script.

### 2.2.3 User Scripting

The last part the job script is the scripting that will be executed by the job. This part of the job script includes all commands required to set up and execute the task your script has been written to do. Any Linux command can be used at this step. This section can be a simple call to an executable or a complex loop which iterates through a series of commands.

Any compute heavy step is preferably should be prefixed by **srun** as the best practice.

Every software program has a unique execution framework. It is the responsibility of the script's author (e.g., you) to know what is required for the software used in your script by reviewing the software's documentation. Regardless of which software your script calls, your script should be written so that the software knows the location of the input and output files as well as the degree of parallelism.

Jobs which touch data-input and data-output files more than once, should make use of `TMPDIR`, a scheduler-provided working space almost 1 TB in size. `TMPDIR` is created when a job starts, and exists on the local disk of the compute node executing your job. Using `TMPDIR` results in faster I/O operations than those to and from shared storage (which is provided over NFS).

An sample job script using `TMPDIR` is available at `/home/n/nul-uge/templateTMPDIR.sh`: the job is instructed to change to `$TMPDIR`, to make the new directory `input`, to copy data from `$SLURM.SUBMIT_DIR/references/` to `input/` (`$SLURM.SUBMIT_DIR` represents the current working directory), to make the new directory `results`, to execute the program (which takes input from `$TMPDIR/input/` and writes output to `$TMPDIR/results/`), and finally to copy the total end results to an existing directory, `processed`, that is located in the current working directory. `TMPDIR` only exists for the duration of the job, though, so it is very important to copy relevant results from it at job's end.

## 2.3 Sample Job Script

Now, let's look at a basic job script, `tcsh.sh` in Figure 3 (you can copy it from our GitHub page or from `/home/n/nul-uge`).

```
#!/encs/bin/tcsh

#SBATCH --job-name=tcsh-test
#SBATCH --mem=1G

sleep 30
module load gurobi/8.1.0
module list
```

Figure 3: Source code for `tcsh.sh`

The first line is the shell declaration (also know as a shebang) and sets the shell to `tcsh`. The lines that begin with `#SBATCH` are directives for the scheduler.

- `-J` (or `--job-name`) sets `tcsh-test` as the job name
- `--chdir` tells the scheduler to execute the job from the current working directory
- `--mem=1GB` requests and assigns 1GB of memory to the job. Jobs *require* the `--mem` option to be set either in the script or on the command line; **if it's missing job submission will be rejected.**

The script then:

- Sleeps on a node for 30 seconds

- Uses the `module` command to load the `gurobi/8.1.0` environment
- Prints the list of loaded modules into a file

The scheduler command, `sbatch`, is used to submit (non-interactive) jobs. From an ssh session on speed-submit, submit this job with `sbatch ./tcsh.sh`. You will see, "Submitted batch job 2653" where 2653 is a job ID assigned. The commands, `squeue` and `sinfo` can be used to look at the status of the cluster: `squeue -l`. You will see something like this:

```
[serguei@speed-submit src] % squeue -l
Thu Oct 19 11:38:54 2023
JOBID PARTITION   NAME       USER      STATE        TIME TIME_LIMI  NODES NODELIST(REASON)
 2641      ps interact   b_user    RUNNING     19:16:09 1-00:00:00      1 speed-07
 2652      ps interact   a_user    RUNNING     41:40 1-00:00:00      1 speed-07
 2654      ps tcsh-tes serguei    RUNNING      0:01 7-00:00:00      1 speed-07
[serguei@speed-submit src] % sinfo
PARTITION AVAIL  TIMELIMIT  NODES  STATE NODELIST
ps*        up 7-00:00:00      14  drain speed-[08-10,12,15-16,20-22,30-32,35-36]
ps*        up 7-00:00:00       1   mix  speed-07
ps*        up 7-00:00:00       7  idle  speed-[11,19,23-24,29,33-34]
pg         up 1-00:00:00       1  drain speed-17
pg         up 1-00:00:00       3  idle  speed-[05,25,27]
pt         up 7-00:00:00       7  idle  speed-[37-43]
pa         up 7-00:00:00       4  idle  speed-[01,03,25,27]
```

Remember that you only have 30 seconds before the job is essentially over, so if you do not see a similar output, either adjust the sleep time in the script, or execute the `sbatch` statement more quickly. The `squeue` output listed above shows you that your job is running on node `speed-07`, that it has a job number of 2654, its time limit of 7 days, etc.

Once the job finishes, there will be a new file in the directory that the job was started from, with the syntax of, `slurm-"job id".out`, so in this example the file is, `slurm-2654.out`. This file represents the standard output (and error, if there is any) of the job in question. If you look at the contents of your newly created file, you will see that it contains the output of the, `module list` command. Important information is often written to this file.

## 2.4 Common Job Management Commands Summary

Here are useful job-management commands:

- `sbatch -A <ACCOUNT> --t <MINUTES> --mem=20G -p <PARTITION> ./<myscript>.sh`: once that your job script is ready, on `speed-submit` you can submit it using this
- `squeue -u <ENCSusername>`: you can check the status of your job(s)
- `squeue`: display cluster status for all users. `-A` shows per account (e.g., `vidpro`, `gipsy`, `speed1`, `ai2`, `aits`, etc.), `-p` per partition (`ps`, `pg`, `pt`, `pa`), and others. `man squeue` for details.
- `squeue --job [job-ID]`: display job information for [job-ID] (said job may be actually running, or waiting in the queue).
- `squeue -las`: displays individual job steps (for debugging easier to see which step failed if you used `srun`).

- `watch -n 1 "sinfo -Nel -pps,pt,pg,pa && squeue -la"`: view `sinfo` information and watch the queue for your job(s).
- `scancel [job-ID]`: cancel job `[job-ID]`.
- `scontrol hold [job-ID]`: hold queued job, `[job-ID]`, from running.
- `scontrol release [job-ID]`: release held job `[job-ID]`.
- `sacct -j [job-ID]`: get job stats. `maxvmem` is one of the more useful stats that you can elect to display as a format option.

```
% sacct -j 2654
```

JobID	JobName	Partition	Account	AllocCPUS	State	ExitCode
2654	tcsh-test	ps	speed1	1	COMPLETED	0:0
2654.batch	batch		speed1	1	COMPLETED	0:0
2654.extern	extern		speed1	1	COMPLETED	0:0

```
% sacct -j 2654 -o jobid,user,account,MaxVMSize,Reason%10,TRESUsageOutMax%30
```

JobID	User	Account	MaxVMSize	Reason	TRESUsageOutMax
2654	serguei	speed1		None	
2654.batch		speed1	296840K		energy=0,fs/disk=1975
2654.extern		speed1	296312K		energy=0,fs/disk=343

See `man sacct` or `sacct -e` for details of the available formatting options. You can define your preferred default format in the `SACCT_FORMAT` environment variable in your `.cshrc` or `.bashrc` files.

- `seff [job-ID]`: reports on the efficiency of a job's cpu and memory utilization. Don't execute it on `RUNNING` jobs (only on completed/finished jobs), efficiency statistics may be misleading.

If you define the following directives in your batch script, you will receive `seff` output in your email when your job is finished.

```
#SBATCH --mail-type=ALL
#SBATCH --mail-user=USER_NAME@encs.concordia.ca
## Replace USER_NAME with your encs username.
```

Output example:

```
Job ID: XXXXX
Cluster: speed
User/Group: user1/user1
State: COMPLETED (exit code 0)
Nodes: 1
Cores per node: 4
CPU Utilized: 00:04:29
CPU Efficiency: 0.35% of 21:32:20 core-walltime
Job Wall-clock time: 05:23:05
Memory Utilized: 2.90 GB
Memory Efficiency: 2.90% of 100.00 GB
```

## 2.5 Advanced sbatch Options

In addition to the basic `sbatch` options presented earlier, there are a few additional options that are generally useful:

- `--mail-type=TYPE`: requests that the scheduler e-mail you when a job changes state. Where `TYPE` is `ALL`, `BEGIN`, `END`, or `FAIL`. Mail is sent to the default address of, "`<ENCSusername>@encs.concordia.ca`", which you can consult via `webmail.encs` via the VPN, on `login.encs` via `alpine` or setup forwarding to `@concordia.ca` address or offsite, unless a different address is supplied (see, `--mail-user`). The report sent when a job ends includes job runtime, as well as the maximum memory value hit (`maxvmem`).
- `--mail-user email@domain.com`: requests that the scheduler use this e-mail notification address, rather than the default (see, `--mail-type`).
- `--export=[ALL | NONE | variables]`: exports environment variable(s) that can be used by the script.
- `-t [min]` or `DAYS-HH:MM:SS`: sets a job runtime of min or HH:MM:SS. Note that if you give a single number, that represents *minutes*, not hours.
- `--depend=[state:job-ID]`: run this job only when job [job-ID] finishes. Held jobs appear in the queue.

The many `sbatch` options available are read with, `man sbatch`. Also note that `sbatch` options can be specified during the job-submission command, and these *override* existing script options (if present). The syntax is, `sbatch [options] PATHTOSCRIPT`, but unlike in the script, the options are specified without the leading `#SBATCH` (e.g., `sbatch -J sub-test --chdir=./ --mem=1G ./tcsh.sh`).

## 2.6 Array Jobs

Array jobs are those that start a batch job or a parallel job multiple times. Each iteration of the job array is called a task and receives a unique job ID. Only supported for batch jobs; submit time < 1 second, compared to repeatedly submitting the same regular job over and over even from a script.

To submit an array job, use the `--array` option of the `sbatch` command as follows:

```
sbatch --array=n-m[:s]] <batch_script>
```

### -t Option Syntax:

- `n`: indicates the start-id.
- `m`: indicates the max-id.
- `s`: indicates the step size.

### Examples:

- `sbatch --array=1-50000 -N1 -i my_in_%a -o my_out_%a array.sh`: submits a job with 50000 elements, `%a` maps to the task-id between 1 and 50K.
- `sbatch --array=10 array.sh`: submits a job with 1 task where the task-id is 10.

- `sbatch --array=1-10 array.sh`: submits a job with 10 tasks numbered consecutively from 1 to 10.
- `sbatch --array=3-15:3 array.sh`: submits a jobs with 5 tasks numbered consecutively with step size 3 (task-ids 3,6,9,12,15).

### Output files for Array Jobs:

The default output and error-files are `slurm-job_id.task_id.out`. This means that Speed creates an output and an error-file for each task generated by the array-job as well as one for the super-ordinate array-job. To alter this behavior use the `-o` and `-e` option of `sbatch`.

For more details about Array Job options, please review the manual pages for `sbatch` by executing the following at the command line on speed-submit `man sbatch`.

## 2.7 Requesting Multiple Cores (i.e., Multithreading Jobs)

For jobs that can take advantage of multiple machine cores, up to 32 cores (per job) can be requested in your script with:

```
#SBATCH -n [#cores for processes]
```

or

```
#SBATCH -n 1
```

```
#SBATCH -c [#cores for threads of a single process]
```

Both `sbatch` and `salloc` support `-n` on the command line, and it should always be used either in the script or on the command line as the default  $n = 1$ . **Do not request more cores than you think will be useful**, as larger-core jobs are more difficult to schedule. On the flip side, though, if you are going to be running a program that scales out to the maximum single-machine core count available, please (please) request 32 cores, to avoid node oversubscription (i.e., to avoid overloading the CPUs).

**Important** note about `--ntasks` or `--ntasks-per-node` (`-n`) talks about processes (usually the ones ran with `srun`). `--cpus-per-task` (`-c`) corresponds to threads per process. Some programs consider them equivalent, some don't. Fluent for example uses `--ntasks-per-node=8` and `--cpus-per-task=1`, some just set `--cpus-per-task=8` and `--ntasks-per-node=1`. If one of them is not 1 then some applications need to be told to use  $n * c$  total cores.

Core count associated with a job appears under, "AllocCPUS", in the, `qacct -j`, output.

```
[serguei@speed-submit src] % squeue -l
Thu Oct 19 20:32:32 2023
JOBID PARTITION NAME USER STATE TIME TIME_LIMI NODES NODELIST(REASON)
2652 ps interact a_user RUNNING 9:35:18 1-00:00:00 1 speed-07
[serguei@speed-submit src] % sacct -j 2652
```

JobID	JobName	Partition	Account	AllocCPUS	State	ExitCode
2652	interacti+	ps	speed1	20	RUNNING	0:0
2652.intera+	interacti+		speed1	20	RUNNING	0:0
2652.extern	extern		speed1	20	RUNNING	0:0
2652.0	gydra_pmi+		speed1	20	COMPLETED	0:0
2652.1	gydra_pmi+		speed1	20	COMPLETED	0:0
2652.2	gydra_pmi+		speed1	20	FAILED	7:0
2652.3	gydra_pmi+		speed1	20	FAILED	7:0

2652.4	gydra_pmi+	speed1	20	COMPLETED	0:0
2652.5	gydra_pmi+	speed1	20	COMPLETED	0:0
2652.6	gydra_pmi+	speed1	20	COMPLETED	0:0
2652.7	gydra_pmi+	speed1	20	COMPLETED	0:0

## 2.8 Interactive Jobs

Job sessions can be interactive, instead of batch (script) based. Such sessions can be useful for testing, debugging, and optimising code and resource requirements, conda or python virtual environments setup, or any likewise preparatory work prior to batch submission.

### 2.8.1 Command Line

To request an interactive job session, use, `salloc [options]`, similarly to a `sbatch` command-line job, e.g.,

```
salloc -J interactive-test --mem=1G -p ps -n 8
```

Inside the allocated `salloc` session you can run shell commands as usual; it is recommended to use `srunch` for the heavy compute steps inside `salloc`. If it is a quick a short job just to compile something, e.g., on a GPU node you can use an interactive `srunch` directly (note no `srunch` can run within `srunch`), e.g., a 1 hour allocation:

For `tcsh`:

```
srunch --pty -n 8 -p pg --gpus=1 --mem=1Gb -t 60 /encs/bin/tcsh
```

For `bash`:

```
srunch --pty -n 8 -p pg --gpus=1 --mem=1Gb -t 60 /encs/bin/bash
```

### 2.8.2 Graphical Applications

If you need to run an on-Speed graphical-based UI application (e.g., MALTLAB, Abaqus CME, etc.), or an IDE (PyCharm, VSCode, Eclipse) to develop and test your job's code interactively you need to enable X11-forwarding from your client machine to speed then to the compute node. To do so:

1. you need to run an X server on your client machine, such as,
  - on Windows: MobaXterm with X turned on, or Xming + PuTTY with X11 forwarding, or XOrg under Cygwin
  - on macOS: XQuartz – use its `xterm` and `ssh -X`
  - on Linux just use `ssh -X speed.encs.concordia.ca`

See <https://www.concordia.ca/ginacody/aits/support/faq/xserver.html> for details.

2. verify your X connection was properly forwarded by printing the `DISPLAY` variable:  

```
echo $DISPLAY
```

 If it has no output, then your X forwarding is not on and you may need to re-login to Speed.
3. Use the `--x11` with `salloc` or `srunch`:  

```
salloc ... --x11=first ...
```



4. Once landed on a compute node, verify `DISPLAY` again.
5. While running under scheduler, create a run-user directory and set the variable `XDGRUNTIME_DIR`.

```
mkdir -p /speed-scratch/$USER/run-dir
setenv XDG_RUNTIME_DIR /speed-scratch/$USER/run-dir
```

6. Launch your graphical application:

module load the required version, then `matlab`, or `abaqus cme`, etc.

Here's an example of starting PyCharm (see Figure 4), of which we made a sample local installation. You can make a similar install under your own directory. If using VSCode, it's currently only supported with the `--no-sandbox` option.

BASH version:

```
bash-3.2$ ssh -X speed (XQuartz xterm, PuTTY or MobaXterm have X11 forwarding too)
serguei@speed's password:
[serguei@speed-submit ~] % echo $DISPLAY
localhost:14.0
[serguei@speed-submit ~] % salloc -p ps --x11=first --mem=4Gb -t 0-06:00
bash-4.4$ echo $DISPLAY
localhost:77.0
bash-4.4$ hostname
speed-01.encs.concordia.ca
bash-4.4$ export XDG_RUNTIME_DIR=/speed-scratch/$USER/run-dir
bash-4.4$ /speed-scratch/nag-public/bin/pycharm.sh
```

TCSH version:

```
ssh -X speed (XQuartz xterm, PuTTY or MobaXterm have X11 forwarding too)
[ speed-submit ] [ /home/c/carlos ] > echo $DISPLAY
localhost:14.0
[ speed-submit ] [ /home/c/carlos ] > cd /speed-scratch/$USER
[ speed-submit ] [ /speed-scratch/carlos ] > echo $DISPLAY
localhost:13.0
[ speed-submit ] [ /speed-scratch/carlos ] > salloc -pps --x11=first --mem=4Gb -t 0-06:00
[ speed-07 ] [ /speed-scratch/carlos ] > echo $DISPLAY
localhost:42.0
[ speed-07 ] [ /speed-scratch/carlos ] > hostname
speed-07.encs.concordia.ca
[ speed-07 ] [ /speed-scratch/carlos ] > setenv XDG_RUNTIME_DIR /speed-scratch/$USER/run-dir
[ speed-07 ] [ /speed-scratch/carlos ] > /speed-scratch/nag-public/bin/pycharm.sh
```

### 2.8.3 Jupyter Notebooks in Singularity

This is an example of running Jupyter notebooks together with Singularity (more on Singularity see Section 2.16). Here we are using one of the OpenISS-derived containers (see Section 2.15.4 as well).

1. Use the `--x11` with `salloc` or `srunc` as described in the above example
2. Load Singularity module `module load singularity/3.10.4/default`
3. Execute this Singularity command on a single line. It's best to save it in a shell script that you could call, since it's long.

```
srunc singularity exec -B $PWD\:/speed-pwd,/speed-scratch/$USER\:/my-speed-scratch,/nettemp \
--env SHELL=/bin/bash --nv /speed-scratch/nag-public/openiss-cuda-conda-jupyter.sif \
/bin/bash -c '/opt/conda/bin/jupyter notebook --no-browser --notebook-dir=/speed-pwd \
--ip="*" --port=8888 --allow-root'
```




Figure 4: PyCharm Starting up on a Speed Node

4. Create an **ssh** tunnel between your computer and the node (**speed-XX**) where Jupyter is running (Using **speed-submit** as a “jump server”) (Preferably: PuTTY, see Figure 5 and Figure 6)

```
ssh -L 8888:speed-XX:8888 YOUR_USER@speed-submit.encs.concordia.ca
```

Don't close the tunnel.

5. Open a browser, and copy your Jupyter's token, in the screenshot example in Figure 7; each time the token will be different, as it printed to you in the terminal.

```
http://localhost:8888/?token=5a52e6c0c7dfc111008a803e5303371ed0462d3d547ac3fb
```

6. Work with your notebook.

#### 2.8.4 Jupyter Labs in Conda and Pytorch

This is an example of Jupyter Labs running in a Conda environment, with Pytorch

- Environment preparation: for the FIRST time:
  1. Go to your speed-scratch directory: `cd /speed-scratch/$USER`
  2. Create a Jupyter (name of your choice) directory: `mkdir -p Jupyter`
  3. Go to Jupyter: `cd Jupyter`
  4. Open an Interactive session: `salloc --mem=50G --gpus=1 -ppg (or -ppt)`

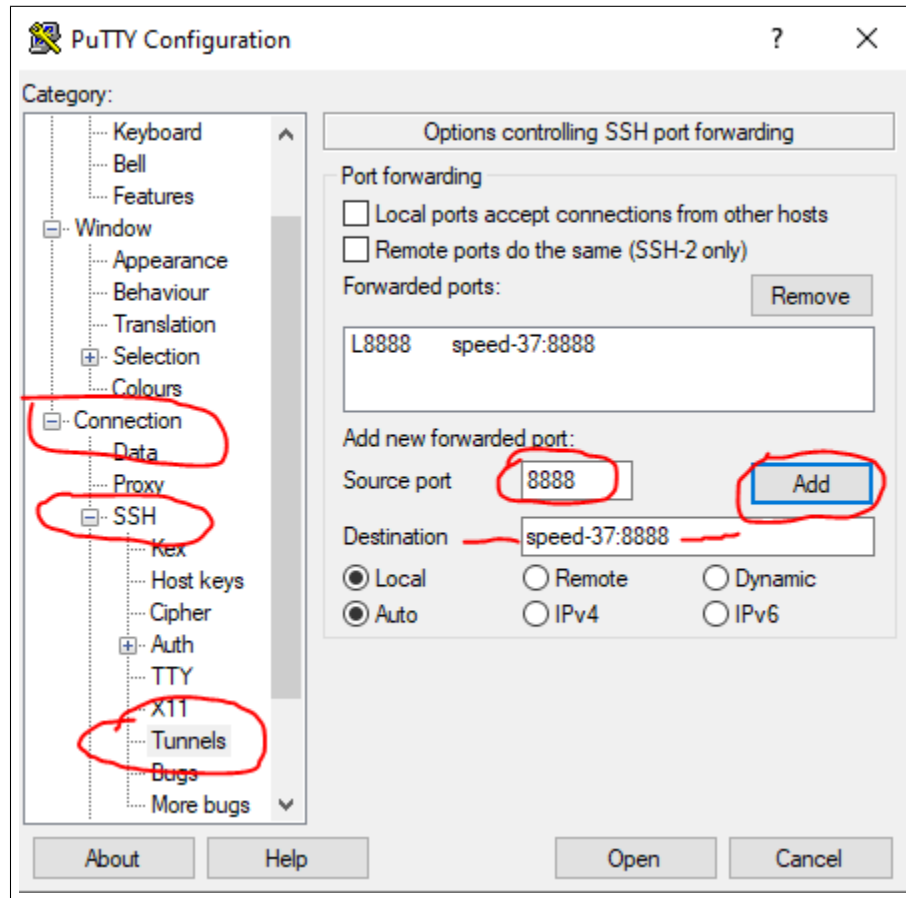


Figure 5: SSH tunnel configuration 1

- Set env. variables, conda environment, jupyter+pytorch installation

```
module load anaconda3/2023.03/default
setenv TMPDIR /speed-scratch/$USER/tmp
setenv TMP /speed-scratch/$USER/tmp
setenv CONDA_PKGS_DIRS /speed-scratch/$USER/Jupyter/pkgs
conda create -p /speed-scratch/$USER/Jupyter/jupyter-env
conda activate /speed-scratch/$USER/Jupyter/jupyter-env
conda install -c conda-forge jupyterlab
pip3 install torch torchvision torchaudio --index-url https://download.pytorch.org/whl/cu118
exit
```

- Running Jupyter Labs, from **speed-submit**:

- Open an Interactive session: `salloc --mem=50G --gpus=1 -ppg (or -ppt)`

```
cd /speed-scratch/$USER/Jupyter
module load anaconda3/2023.03/default
setenv TMPDIR /speed-scratch/$USER/tmp
setenv TMP /speed-scratch/$USER/tmp
setenv CONDA_PKGS_DIRS /speed-scratch/$USER/Jupyter/pkgs
conda activate /speed-scratch/$USER/Jupyter/jupyter-env
jupyter lab --no-browser --notebook-dir=$PWD --ip="*" --port=8888 --port-retries=50
```

- Verify which port the system has assigned to Jupyter: `http://localhost:XXXX/lab?token=`

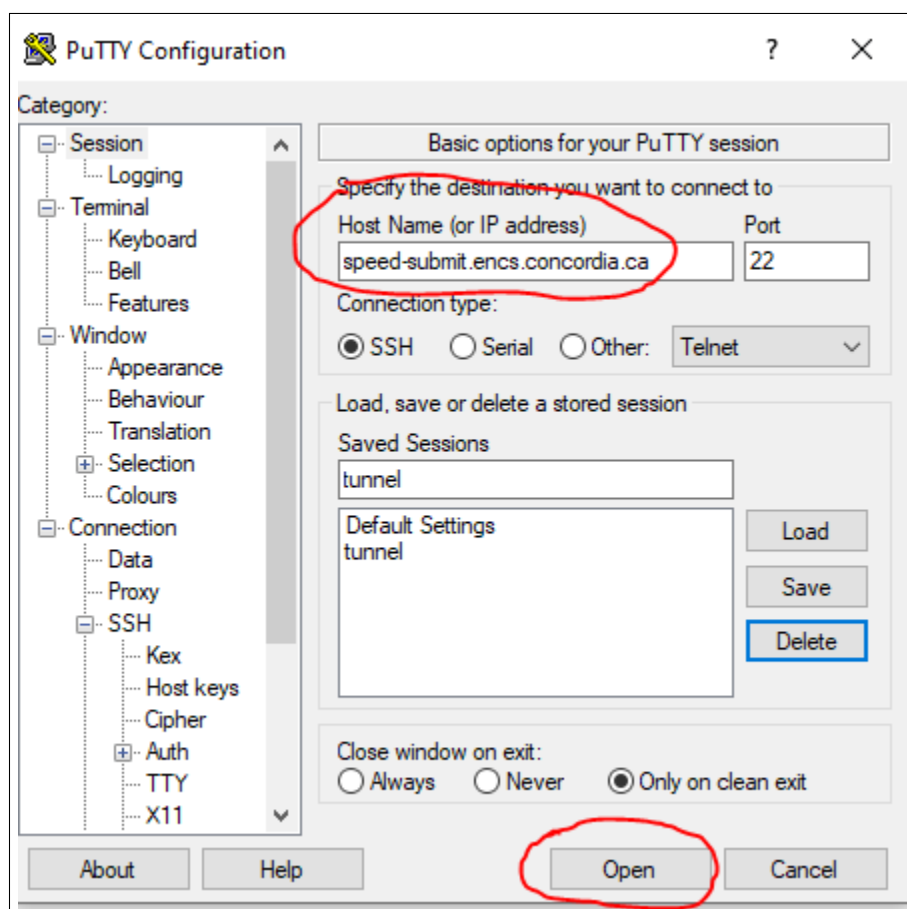


Figure 6: SSH tunnel configuration 2




Figure 7: Jupyter running on a Speed node

3. SSH Tunnel creation: similar to Jupyter in Singularity, see Section 2.8.3
4. Open a browser and type: `localhost:XXXX` (port assigned)

### 2.8.5 VScode

This is an example of running VScode, it's similar to Jupyter notebooks, but it doesn't use containers. This a Web version, it exists the local(workstation)-remote(speed-node) version too, but it is for Advanced users (no support, execute it at your own risk).

- Environment preparation: for the FIRST time:
  1. Go to your speed-scratch directory: `cd /speed-scratch/$USER`
  2. Create a vscode directory: `mkdir vscode`
  3. Go to vscode: `cd vscode`
  4. Create home and projects: `mkdir {home,projects}`
  5. Create this directory: `mkdir -p /speed-scratch/$USER/run-user`
- Running VScode
  1. Go to your vscode directory: `cd /speed-scratch/$USER/vscode`
  2. Open interactive session: `salloc --mem=10Gb --constraint=el9`
  3. Set environment variable: `setenv XDG_RUNTIME_DIR /speed-scratch/$USER/run-user`
  4. Run VScode, change the port if needed.
 

```
/speed-scratch/nag-public/code-server-4.22.1/bin/code-server --user-data-dir=$PWD\projects \
--config=$PWD\home/.config/code-server/config.yaml --bind-addr="0.0.0.0:8080" $PWD\projects
```
  5. SSH Tunnel creation: similar to Jupyter, see Section 2.8.3
  6. Open a browser and type: `localhost:8080`
  7. If the browser asks for password:
 

```
cat /speed-scratch/$USER/vscode/home/.config/code-server/config.yaml
```

## 2.9 Scheduler Environment Variables

The scheduler presents a number of environment variables that can be used in your jobs. You can invoke `env` or `printenv` in your job to know what those are (most begin with the prefix `SLURM`). Some of the more useful ones are:

- `$TMPDIR` – the path to the job's temporary space on the node. It *only* exists for the duration of the job, so if data in the temporary space are important, they absolutely need to be accessed before the job terminates.
- `$SLURM_SUBMIT_DIR` – the path to the job's working directory (likely an NFS-mounted path). If, `--chdir`, was stipulated, that path is taken; if not, the path defaults to your home directory.
- `$SLURM_JOBID` – your current jobs ID, useful for some manipulation and reporting.
- `$SLURM_JOB_NODELIST`=nodes participating in your job.




Figure 8: VScode running on a Speed node

- `$SLURM_ARRAY_TASK_ID`=for array jobs (see Section 2.6).
- See a more complete list here:
  - [https://slurm.schedmd.com/srun.html#SECTION\\_INPUT-ENVIRONMENT-VARIABLES](https://slurm.schedmd.com/srun.html#SECTION_INPUT-ENVIRONMENT-VARIABLES)
  - [https://slurm.schedmd.com/srun.html#SECTION\\_OUTPUT-ENVIRONMENT-VARIABLES](https://slurm.schedmd.com/srun.html#SECTION_OUTPUT-ENVIRONMENT-VARIABLES)

In Figure 9 is a sample script, using some of these.

```
#!/encs/bin/tcsh

#SBATCH --job-name=tmpdir      ## Give the job a name
#SBATCH --mail-type=ALL       ## Receive all email type notifications
#SBATCH --mail-user=YOUR_USER_NAME@encs.concordia.ca
#SBATCH --chdir=./            ## Use current directory as working directory
#SBATCH --nodes=1
#SBATCH --ntasks=1
#SBATCH --cpus-per-task=8     ## Request 8 cores
#SBATCH --mem=32G             ## Assign 32G memory per node

cd $TMPDIR
mkdir input
rsync -av $SLURM_SUBMIT_DIR/references/ input/
mkdir results
srun STAR --inFiles $TMPDIR/input --parallel $SRUN_CPUS_PER_TASK --outFiles $TMPDIR/results
rsync -av $TMPDIR/results/ $SLURM_SUBMIT_DIR/processed/
```

Figure 9: Source code for `tmpdir.sh`

## 2.10 SSH Keys For MPI

Some programs effect their parallel processing via MPI (which is a communication protocol). An example of such software is Fluent. MPI needs to have ‘passwordless login’ set up, which

means SSH keys. In your NFS-mounted home directory:

- `cd .ssh`
- `ssh-keygen -t ed25519` (default location; blank passphrase)
- `cat id_ed25519.pub >> authorized_keys` (if the `authorized_keys` file already exists)  
*OR* `cat id_ed25519.pub > authorized_keys` (if does not)
- Set file permissions of `authorized_keys` to 600; of your NFS-mounted home to 700 (note that you likely will not have to do anything here, as most people will have those permissions by default).

## 2.11 Creating Virtual Environments

The following documentation is specific to the **Speed** HPC Facility at the Gina Cody School of Engineering and Computer Science. Virtual environments typically instantiated via Conda or Python. Another option is Singularity detailed in Section 2.16. Usually, virtual environments are created once during an interactive session before submitting a batch job to the scheduler. The job script submitted to the scheduler is then written to (1) activate the virtual environment, (2) use it, and (3) close it at the end of the job.

### 2.11.1 Anaconda

Request an interactive session in the queue you wish to submit your jobs to (e.g., `salloc -p pg -gpus=1` for GPU jobs). Once your interactive has started, create an anaconda environment in your speed-scratch directory by using the `prefix` option when executing `conda create`. For example, to create an anaconda environment for `a_user`, execute the following at the command line:

```
module load anaconda3/2023.03/default
conda create --prefix /speed-scratch/a_user/myconda
```

**Note:** Without the `prefix` option, the `conda create` command creates the environment in `a_user`'s home directory by default.

**List Environments.** To view your conda environments, type: `conda info --envs`

```
# conda environments:
#
base                  *  /encs/pkg/anaconda3-2023.03/root
                        /speed-scratch/a_user/myconda
```

**Activate an Environment.** Activate the environment `speedscratcha_usermyconda` as follows

```
conda activate /speed-scratch/a_user/myconda
```

After activating your environment, add `pip` to your environment by using

```
conda install pip
```

This will install `pip` and `pip`'s dependencies, including `python`, into the environment.

- A consolidated example using Conda:

```
salloc -p pg --gpus=1 --mem=10GB -A <slurm account name>
cd /speed-scratch/$USER
module load python/3.11.0/default
conda create -p /speed-scratch/$USER/pytorch-env
conda activate /speed-scratch/$USER/pytorch-env
conda install python=3.11.0
pip3 install torch torchvision torchaudio --index-url \
    https://download.pytorch.org/whl/cu117
....
conda deactivate
exit # end the salloc session
```

- No Space left error: Read our Github [HERE](#)

**Important Note:** `pip` (and `pip3`) are used to install modules from the python distribution while `conda install` installs modules from anaconda's repository.

### 2.11.2 Python

Setting up a Python virtual environment is fairly straightforward. The first step is to request an interactive session in the queue you wish to submit your jobs to.

We have a simple example that use a Python virtual environment:

- Using Python Venv

```
salloc -p pg --gpus=1 --mem=10GB -A <slurm account name>
cd /speed-scratch/$USER
module load python/3.9.1/default
mkdir -p /speed-scratch/$USER/tmp
setenv TMPDIR /speed-scratch/$USER/tmp
setenv TMP /speed-scratch/$USER/tmp
python -m venv $TMPDIR/testenv (testenv=name of the virtualEnv)
source /speed-scratch/$USER/tmp/testenv/bin/activate.csh
pip install modules...
deactivate
exit
```

- See, e.g., `gurobi-with-python.sh`

**Important Note:** partition `ps` is used for CPU jobs, partitions `pg`, `pt` are used for GPU jobs, no need to use `--gpus=` when preparing environments for CPU jobs.



```
#!/encs/bin/tcsh

#SBATCH --job-name=flu10000    ## Give the job a name
#SBATCH --mail-type=ALL       ## Receive all email type notifications
#SBATCH --mail-user=YOUR_USER_NAME@encs.concordia.ca
#SBATCH --chdir=./            ## Use current directory as working directory
#SBATCH --nodes=1             ## Number of nodes to run on
#SBATCH --ntasks-per-node=32  ## Number of cores
#SBATCH --cpus-per-task=1     ## Number of MPI threads
#SBATCH --mem=160G            ## Assign 160G memory per node

date

module avail ansys

module load ansys/19.2/default
cd $TMPDIR

set FLUENTNODES = "`scontrol show hostnames`"
set FLUENTNODES = `echo $FLUENTNODES | tr ' ' ','`

date

srun fluent 3ddp \
    -g -t$SLURM_NTASKS \
    -g-cnf=$FLUENTNODES \
    -i $SLURM_SUBMIT_DIR/fluentsdata/info.jou > call.txt

date

srun rsync -av $TMPDIR/ $SLURM_SUBMIT_DIR/fluentsparallel/

date
```

Figure 10: Source code for fluent.sh

## 2.12 Example Job Script: Fluent

The job script in Figure 10 runs Fluent in parallel over 32 cores. Of note, we have requested e-mail notifications (`--mail-type`), are defining the parallel environment for, `fluent`, with, `-t$SLURM_NTASKS` and `-g-cnf=$FLUENTNODES` (**very important**), and are setting `$TMPDIR` as the in-job location for the “moment” `rfile.out` file (in-job, because the last line of the script copies everything from `$TMPDIR` to a directory in the user’s NFS-mounted home). Job progress can be monitored by examining the standard-out file (e.g., `slurm-249.out`), and/or by examining the “moment” file in `/disk/nobackup/<yourjob>` (hint: it starts with your job-ID) on the node running the job. **Caveat:** take care with journal-file file paths.

## 2.13 Example Job: efficientdet

The following steps describing how to create an `efficientdet` environment on *Speed*, were submitted by a member of Dr. Amer’s research group.

- Enter your ENCS user account’s speed-scratch directory  
`cd /speed-scratch/<encs_username>`

- Next
  - load python module `load python/3.8.3`
  - create virtual environment `python3 -m venv <env_name>`
  - activate virtual environment `source <env_name>/bin/activate.csh`
  - install DL packages for Efficientdet

```

pip install tensorflow==2.7.0
pip install lxml>=4.6.1
pip install absl-py>=0.10.0
pip install matplotlib>=3.0.3
pip install numpy>=1.19.4
pip install Pillow>=6.0.0
pip install PyYAML>=5.1
pip install six>=1.15.0
pip install tensorflow-addons>=0.12
pip install tensorflow-hub>=0.11
pip install neural-structured-learning>=1.3.1
pip install tensorflow-model-optimization>=0.5
pip install Cython>=0.29.13
pip install git+https://github.com/cocodataset/cocoapi.git#subdirectory=PythonAPI

```

## 2.14 Java Jobs

Jobs that call `java` have a memory overhead, which needs to be taken into account when assigning a value to `--mem`. Even the most basic `java -Xmx1G -version`, will need to have, `--mem=5G`, with the 4-GB difference representing the memory overhead. Note that this memory overhead grows proportionally with the value of `-Xmx`. To give you an idea, when `-Xmx` has a value of 100G, `--mem` has to be at least 106G; for 200G, at least 211G; for 300G, at least 314G.

## 2.15 Scheduling On The GPU Nodes

The primary cluster has two GPU nodes, each with six Tesla (CUDA-compatible) P6 cards: each card has 2048 cores and 16GB of RAM. Though note that the P6 is mainly a single-precision card, so unless you need the GPU double precision, double-precision calculations will be faster on a CPU node.

Job scripts for the GPU queue differ in that they need this statement, which attaches either a single GPU, or, two GPUs, to the job:

```
#SBATCH --gpus=[1|2]
```

Once that your job script is ready, you can submit it to the GPU partition (queue) with:

```
sbatch -p pg ./<myscript>.sh
```

And you can query `nvidia-smi` on the node that is running your job with:

```
ssh <username>@speed[-05|-17|37-43] nvidia-smi
```

Status of the GPU queue can be queried with:

```
sinfo -p pg --long --Node
```

**Very important note** regarding TensorFlow and PyTorch: if you are planning to run TensorFlow and/or PyTorch multi-GPU jobs, **do not** use the `tf.distribute` and/or `torch.nn.DataParallel` functions on **speed-01,05,17**, as they will crash the compute node (100% certainty). This appears to be the current hardware's architecture's defect. The workaround is to either manually effect GPU parallelisation (TensorFlow has an example on how to do this), or to run on a single GPU.

## Important

Users without permission to use the GPU nodes can submit jobs to the `pg` partition, but those jobs will hang and never run. Their availability is seen with:

```
[serguei@speed-submit src] % sinfo -p pg --long --Node
Thu Oct 19 22:31:04 2023
NODELIST  NODES PARTITION  STATE CPUS   S:C:T MEMORY TMP_DISK WEIGHT AVAIL_FE REASON
speed-05   1      pg          idle  32    2:16:1 515490      0      1    gpu16 none
speed-17   1      pg          drained 32    2:16:1 515490      0      1    gpu16 UGE
speed-25   1      pg          idle  32    2:16:1 257458      0      1    gpu32 none
speed-27   1      pg          idle  32    2:16:1 257458      0      1    gpu32 none
[serguei@speed-submit src] % sinfo -p pt --long --Node
Thu Oct 19 22:32:39 2023
NODELIST  NODES PARTITION  STATE CPUS   S:C:T MEMORY TMP_DISK WEIGHT AVAIL_FE REASON
speed-37   1      pt          idle  256    2:64:2 980275      0      1    gpu20,mi none
speed-38   1      pt          idle  256    2:64:2 980275      0      1    gpu20,mi none
speed-39   1      pt          idle  256    2:64:2 980275      0      1    gpu20,mi none
speed-40   1      pt          idle  256    2:64:2 980275      0      1    gpu20,mi none
speed-41   1      pt          idle  256    2:64:2 980275      0      1    gpu20,mi none
speed-42   1      pt          idle  256    2:64:2 980275      0      1    gpu20,mi none
speed-43   1      pt          idle  256    2:64:2 980275      0      1    gpu20,mi none
```

This status demonstrates that most are available (i.e., have not been requested as resources). To specifically request a GPU node, add, `--gpus=[#GPUs]`, to your `sbatch` (statement/script) or `salloc` (statement) request. For example, `sbatch -t 10 --mem=1G --gpus=1 -p pg ./tcsh.sh`. You will see that this job has been assigned to one of the GPU nodes.

```
[serguei@speed-submit src] % squeue -p pg -o "%15N %.6D %7P %.11T %.4c %.8z %.6m %.8d %.6w %.8f %20G %20E"
NODELIST  NODES PARTITI  STATE MIN_   S:C:T MIN_ME MIN_TMP_  WCKEY FEATURES GROUP DEPENDENCY
speed-05   1      pg          RUNNING 1    *:*:*      1G      0 (null) (null) 11929 (null)
[serguei@speed-submit src] % sinfo -p pg -o "%15N %.6D %7P %.11T %.4c %.8z %.6m %.8d %.6w %.8f %20G %20E"
NODELIST  NODES PARTITI  STATE CPUS   S:C:T MEMORY TMP_DISK WEIGHT AVAIL_FE GRES      REASON
speed-17   1      pg          drained 32    2:16:1 515490      0      1    gpu16 gpu:6    UGE
speed-05   1      pg          mixed  32    2:16:1 515490      0      1    gpu16 gpu:6    none
speed-[25,27] 2      pg          idle  32    2:16:1 257458      0      1    gpu32 gpu:2    none
```

### 2.15.1 P6 on Multi-GPU, Multi-Node

As described lines above, P6 cards are not compatible with `Distribute` and `DataParallel` functions (`Pytorch`, `Tensorflow`) when running on Multi-GPUs. One workaround is to run the job in Multi-node, single GPU per node; per example:

```
#SBATCH --nodes=2
#SBATCH --gpus-per-node=1
```

On P6 nodes: `speed-05`, `speed-17`, `speed-01`

The example: `pytorch-multinode-multigpu.sh` illustrates a job for training on Multi-nodes, Multi-GPUs

### 2.15.2 CUDA

When calling CUDA within job scripts, it is important to create a link to the desired CUDA libraries and set the runtime link path to the same libraries. For example, to use the `cuda-11.5` libraries, specify the following in your `Makefile`.

```
-L/encs/pkg/cuda-11.5/root/lib64 -Wl,-rpath,/encs/pkg/cuda-11.5/root/lib64
```

In your job script, specify the version of `gcc` to use prior to calling `cuda`. For example:  
`module load gcc/8.4` or `module load gcc/9.3`

### 2.15.3 Special Notes for sending CUDA jobs to the GPU Queue

Interactive jobs (Section 2.8) must be submitted to the **GPU partition** in order to compile and link. We have several versions of CUDA installed in:

```
/encs/pkg/cuda-11.5/root/  
/encs/pkg/cuda-10.2/root/  
/encs/pkg/cuda-9.2/root
```

For CUDA to compile properly for the GPU partition, edit your `Makefile` replacing `usrlocalcuda` with one of the above.

### 2.15.4 OpenISS Examples

These represent more comprehensive research-like examples of jobs for computer vision and other tasks with a lot longer runtime (a subject to the number of epochs and other parameters) derive from the actual research works of students and their theses. These jobs require the use of CUDA and GPUs. These examples are available as “native” jobs on Speed and as Singularity containers.

**OpenISS and REID** The example `openiss-reid-speed.sh` illustrates a job for a computer-vision based person re-identification (e.g., motion capture-based tracking for stage performance) part of the OpenISS project by Haotao Lai [10] using TensorFlow and Keras. The fork of the original repo [12] adjusted to to run on Speed is here:

- <https://github.com/NAG-DevOps/openiss-reid-tfk>

and its detailed description on how to run it on Speed is in the README:

- <https://github.com/NAG-DevOps/speed-hpc/tree/master/src#openiss-reid-tfk>

**OpenISS and YOLOv3** The related code using YOLOv3 framework is in the the fork of the original repo [11] adjusted to to run on Speed is here:

- <https://github.com/NAG-DevOps/openiss-yolov3>

Its example job scripts can run on both CPUs and GPUs, as well as interactively using TensorFlow:

- Interactive mode: `openiss-yolo-interactive.sh`

- CPU-based job: `openiss-yolo-cpu.sh`
- GPU-based job: `openiss-yolo-gpu.sh`

The detailed description on how to run these on Speed is in the README at:

- <https://github.com/NAG-DevOps/speed-hpc/tree/master/src#openiss-yolov3>

## 2.16 Singularity Containers

If the `/encs` software tree does not have a required software instantaneously available, another option is to run Singularity containers. We run EL7 flavor of Linux, and if some projects require Ubuntu or other distributions, there is a possibility to run that software as a container, including the ones translated from Docker.

The example `lambdal-singularity.sh` showcases an immediate use of a container built for the Ubuntu-based LambdaLabs software stack, originally built as a Docker image then pulled in as a Singularity container that is immediately available for use as that job example illustrates. The source material used for the docker image was our fork of their official repo: <https://github.com/NAG-DevOps/lambda-stack-dockerfiles>

NOTE: It is important if you make your own containers or pull from DockerHub, use your `/speed-scratch/$USER` directory as these images may easily consume gigs of space in your home directory and you'd run out of quota there very fast.

TIP: To check for your quota, and the corresponding commands to find big files, see: <https://www.concordia.ca/ginacody/aits/encs-data-storage.html>

We likewise built equivalent OpenISS (Section 2.15.4) containers from their Docker counterparts as they were used for teaching and research [14]. The images from <https://github.com/NAG-DevOps/openiss-dockerfiles> and their DockerHub equivalents <https://hub.docker.com/u/openiss> are found in the same public directory on `/speed-scratch/nag-public` as the LambdaLabs Singularity image. They all have `.sif` extension. Some of them can be ran in both batch or interactive mode, some make more sense to run interactively. They cover some basics with CUDA, OpenGL rendering, and computer vision tasks as examples from the OpenISS library and other libraries, including the base images that use different distros. We also include Jupyter notebook example with Conda support.

`/speed-scratch/nag-public:`

```
openiss-cuda-conda-jupyter.sif
openiss-cuda-devicequery.sif
openiss-opengl-base.sif
openiss-opengl-cubes.sif
openiss-opengl-triangle.sif
openiss-reid.sif
openiss-xeyes.sif
```

The currently recommended version of Singularity is `singularity/3.10.4/default`.

This section comprises an introduction to working with Singularity, its containers, and what can and cannot be done with Singularity on the ENCS infrastructure. It is not intended to be an exhaustive presentation of Singularity: the program's authors do a good job of that here: <https://www.sylabs.io/docs/>. It also assumes that you have successfully installed Singularity on a user-managed/personal system (see next paragraph as to why).

Singularity containers are essentially either built from an existing container, or are built from scratch. Building from scratch requires a recipe file (think of like a Dockerfile), and the operation *must* be effected as root. You will not have root on the ENCS infrastructure, so any built-from-scratch containers must be created on a user-managed/personal system. Root-level permissions are also required (in some cases, essential; in others, for proper build functionality) for building from an existing container. Three types of Singularity containers can be built: file-system; sandbox; squashfs. The first two are “writable” (meaning that changes can persist after the Singularity session ends). File-system containers are built around the ext3 file system, and are a read-write “file”, sandbox containers are essentially a directory in an existing read-write space, and squashfs containers are a read-only compressed “file”. Note that file-system containers *cannot* be resized once built.

Note that the default build is a squashfs one. Also note what Singularity’s authors have to say about the builds, “A common workflow is to use the “sandbox” mode for development of the container, and then build it as a default (squashfs) Singularity image when done.” File-system containers are considered to be, “legacy”, at this point in time. When built, a *very small* overhead is allotted to a file-system container (think, MB), and that *cannot* be changed.

Probably for the most of your workflows you might find there is a Docker container exists for your tasks, in this case you can use the docker pull function of Singularity as a part of you virtual environment setup as an interactive job allocation:

```
salloc --gpus=1 -n8 --mem=4Gb -t60
cd /speed-scratch/$USER/
singularity pull openiss-cuda-devicequery.sif docker://openiss/openiss-cuda-devicequery
INFO:    Converting OCI blobs to SIF format
INFO:    Starting build...
```

This method can be used for converting Docker containers directly on Speed. On GPU nodes make sure to pass on the `--nv` flag to Singularity, so its containers could access the GPUs. See the linked example.

## 3 Conclusion

The cluster is, “first come, first served”, until it fills, and then job position in the queue is based upon past usage. The scheduler does attempt to fill gaps, though, so sometimes a single-core job of lower priority will schedule before a multi-core job of higher priority, for example.

### 3.1 Important Limitations

- New users are restricted to a total of 32 cores: write to `rt-ex-hpc@encs.concordia.ca` if you need more temporarily (192 is the maximum, or, 6 jobs of 32 cores each).
- Batch job sessions are a maximum of one week in length (only 24 hours, though, for interactive jobs, see Section 2.8).
- Scripts can live in your NFS-provided home, but any substantial data need to be in your cluster-specific directory (located at `/speed-scratch/<ENCSusername>/`).

NFS is great for acute activity, but is not ideal for chronic activity. Any data that a job will read more than once should be copied at the start to the scratch disk of a compute node using `$TMPDIR` (and, perhaps, `$SLURM_SUBMIT_DIR`), any intermediary job data should be

produced in `$TMPDIR`, and once a job is near to finishing, those data should be copied to your NFS-mounted home (or other NFS-mounted space) from `$TMPDIR` (to, perhaps, `$SLURM_SUBMIT_DIR`). In other words, IO-intensive operations should be effected locally whenever possible, saving network activity for the start and end of jobs.

- Your current resource allocation is based upon past usage, which is an amalgamation of approximately one week’s worth of past wallclock (i.e., time spent on the node(s)) and compute activity (on the node(s)).
- Jobs should NEVER be run outside of the province of the scheduler. Repeat offenders risk loss of cluster access.

### 3.2 Tips/Tricks

- Files/scripts must have Linux line breaks in them (not Windows ones). Use `file` command to verify; and `dos2unix` command to convert.
- Use `rsync`, not `scp`, when copying or moving large amounts of data.
- Before moving a large amount of files between NFS-mounted storage and the cluster, `tar` up the files you plan to move first.
- If you intend to use a different shell (e.g., `bash` [22]), you will need to change the shell declaration in your script(s).
- **Try to request resources that closely match what your job will use: requesting many more cores or much more memory than will be needed makes a job more difficult to schedule when resources are scarce.**
- E-mail, `rt-ex-hpc AT encs.concordia.ca`, with any concerns/questions.

### 3.3 Use Cases

- HPC Committee’s initial batch about 6 students (end of 2019):
  - 10000 iterations job in Fluent finished in < 26 hours vs. 46 hours in Calcul Quebec
- NAG’s MAC spoofer analyzer [18, 17], such as <https://github.com/smokhov/at-sm/tree/master/examples/flucid>
  - compilation of forensic computing reasoning cases about false or true positives of hardware address spoofing in the labs
- S4 LAB/GIPSY R&D Group’s:
  - MARFCAT and MARFPCAT (OSS signal processing and machine learning tools for vulnerable and weak code analysis and network packet capture analysis) [20, 15, 6]
  - Web service data conversion and analysis
  - Forensic Lucid encoders (translation of large log data into Forensic Lucid [16] for forensic analysis)
  - Genomic alignment exercises

- Serguei Mokhov, Jonathan Llewellyn, Carlos Alarcon Meza, Tariq Daradkeh, and Gillian Roper. The use of containers in OpenGL, ML and HPC for teaching and research support. In *ACM SIGGRAPH 2023 Posters*, SIGGRAPH '23, New York, NY, USA, 2023. ACM. <https://doi.org/10.1145/3588028.3603676>
- Goutam Yelluru Gopal and Maria Amer. Separable self and mixed attention transformers for efficient object tracking. In *IEEE/CVF Winter Conference on Applications of Computer Vision (WACV)*, Waikoloa, Hawaii, January 2024. <https://arxiv.org/abs/2309.03979> and <https://github.com/goutamylg/SMAT>
- Goutam Yelluru Gopal and Maria Amer. Mobile vision transformer-based visual object tracking. In *34th British Machine Vision Conference (BMVC)*, Aberdeen, UK, November 2023. <https://arxiv.org/abs/2309.05829> and <https://github.com/goutamylg/MVT>
- Belkacem Belabes and Marius Paraschivoiu. CFD modeling of vertical-axis wind turbine wake interaction. *Transactions of the Canadian Society for Mechanical Engineering*, pages 1–10, 2023. <https://doi.org/10.1139/tcsme-2022-0149>
- Belkacem Belabes and Marius Paraschivoiu. CFD study of the aerodynamic performance of a vertical axis wind turbine in the wake of another turbine. In *Proceedings of the CSME International Congress*, 2022. <https://doi.org/10.7939/r3-rker-1746>
- Belkacem Belabes and Marius Paraschivoiu. Numerical study of the effect of turbulence intensity on VAWT performance. *Energy*, 233:121139, 2021. <https://doi.org/10.1016/j.energy.2021.121139>
- Parna Niksirat, Adriana Daca, and Krzysztof Skonieczny. The effects of reduced-gravity on planetary rover mobility. *International Journal of Robotics Research*, 39(7):797–811, 2020. <https://doi.org/10.1177/0278364920913945>
- The work “Haotao Lai. An OpenISS framework specialization for deep learning-based person re-identification. Master’s thesis, Department of Computer Science and Software Engineering, Concordia University, Montreal, Canada, August 2019. <https://spectrum.library.concordia.ca/id/eprint/985788/>” using TensorFlow and Keras on OpenISS adjusted to run on Speed based on the repositories:
  - Haotao Lai et al. Openiss person re-identification baseline v0.1.1, June 2021. <https://github.com/OpenISS/openiss-reid-tfk> and
  - Haotao Lai et al. OpenISS keras-yolo3 v0.1.0, June 2021. <https://github.com/OpenISS/openiss-yolov3>

and theirs forks by the team.

## A History

### A.1 Acknowledgments

- The first 6 (to 6.5) versions of this manual and early UGE job script samples, Singularity testing and user support were produced/done by Dr. Scott Bunnell during his time at Concordia as a part of the NAG/HPC group. We thank him for his contributions.
- The HTML version with devcontainer support was contributed by Anh H Nguyen.



- Dr. Tariq Daradkeh, was our IT Instructional Specialist August 2022 to September 2023; working on the scheduler, scheduling research, end user support, and integration of examples, such as YOLOv3 in Section 2.15.4 other tasks. We have a continued collaboration on HPC/scheduling research.

## A.2 Migration from UGE to SLURM

For long term users who started off with Grid Engine here are some resources to make a transition and mapping to the job submission process.

- Queues are called “partitions” in SLURM. Our mapping from the GE queues to SLURM partitions is as follows:

```
GE  => SLURM
s.q   ps
g.q   pg
a.q   pa
```

We also have a new partition `pt` that covers SPEED2 nodes, which previously did not exist.

- Commands and command options mappings are found in Figure 11 from <https://slurm.schedmd.com/rosetta.pdf>  
<https://slurm.schedmd.com/pdfs/summary.pdf>  
 Other related helpful resources from similar organizations who either used SLURM for awhile or also transitioned to it:  
[https://docs.alliancecan.ca/wiki/Running\\_jobs](https://docs.alliancecan.ca/wiki/Running_jobs)  
[https://www.depts.ttu.edu/hpcc/userguides/general\\_guides/Conversion\\_Table\\_1.pdf](https://www.depts.ttu.edu/hpcc/userguides/general_guides/Conversion_Table_1.pdf)  
<https://docs.mpcdf.mpg.de/doc/computing/clusters/aux/migration-from-sge-to-slurm>
- **NOTE:** If you have used UGE commands in the past you probably still have these lines there; **they should now be removed**, as they have no use in SLURM and will start giving “command not found” errors on login when the software is removed:

csh/tcsh: Sample `.tcshrc` file:

```
# Speed environment set up
if ($HOSTNAME == speed-submit.encs.concordia.ca) then
    source /local/pkg/uge-8.6.3/root/default/common/settings.csh
endif
```

Bourne shell/bash: Sample `.bashrc` file:

```
# Speed environment set up
if [ $HOSTNAME = "speed-submit.encs.concordia.ca" ]; then
    . /local/pkg/uge-8.6.3/root/default/common/settings.sh
    printenv ORGANIZATION | grep -qw ENCS || . /encs/Share/bash/profile
fi
```

Note that you will need to either log out and back in, or execute a new shell, for the environment changes in the updated `.tcshrc` or `.bashrc` file to be applied (**important**).

User Commands	PBS/Torque	Slurm	LSF	SGE
Job submission	qsub [script_file]	sbatch [script_file]	bsub [script_file]	qsub [script_file]
Job deletion	qdel [job_id]	scancel [job_id]	bkill [job_id]	qdel [job_id]
Job status (by job)	qstat [job_id]	squeue [job_id]	bjobs [job_id]	qstat -u '*' [-j job_id]
Job status (by user)	qstat -u [user_name]	squeue -u [user_name]	bjobs -u [user_name]	qstat [-u user_name]
Job hold	qhold [job_id]	scontrol hold [job_id]	bstop [job_id]	qhold [job_id]
Job release	qrls [job_id]	scontrol release [job_id]	brresume [job_id]	qrls [job_id]
Queue list	qstat -Q	squeue	bqueues	qconf -sq
Node list	pbsnodes -l	sinfo -N OR scontrol show nodes	bhosts	qhost
Cluster status	qstat -a	sinfo	bqueues	qhost -q
GUI	xpbsmon	sview	xlsf OR xlsbatch	qmon
Environment	PBS/Torque	Slurm	LSF	SGE
Job ID	\$PBS_JOBID	\$SLURM_JOBID	\$LSB_JOBID	\$JOB_ID
Submit Directory	\$PBS_O_WORKDIR	\$SLURM_SUBMIT_DIR	\$LSB_SUBCWD	\$SGE_O_WORKDIR
Submit Host	\$PBS_O_HOST	\$SLURM_SUBMIT_HOST	\$LSB_SUB_HOST	\$SGE_O_HOST
Node List	\$PBS_NODEFILE	\$SLURM_JOB_NODELIST	\$LSB_HOSTS/LSB_MCPU_HOST	\$PE_HOSTFILE
Job Array Index	\$PBS_ARRAYID	\$SLURM_ARRAY_TASK_ID	\$LSB_JOBINDEX	\$SGE_TASK_ID
Job Specification	PBS/Torque	Slurm	LSF	SGE
Script directive	#PBS	#SBATCH	#BS	#\$
Queue	-q [queue]	-p [queue]	-q [queue]	-q [queue]
Node Count	-l nodes=[count] -l ppn=[count] OR -l mppwidth=[PE_count]	-N [min[-max]] -n [count]	-n [count]	N/A
CPU Count				-pe [PE] [count]
Wall Clock Limit	-l walltime=[hh:mm:ss]	-t [min] OR -t [days-hh:mm:ss]	-W [hh:mm:ss]	-l h_rt=[seconds]
Standard Output File	-o [file_name]	-o [file_name]	-o [file_name]	-o [file_name]
Standard Error File	-e [file_name]	-e [file_name]	-e [file_name]	-e [file_name]
Combine stdout/err	-j oe (both to stdout) OR -j eo (both to stderr)	(use -o without -e)	(use -o without -e)	-j yes
Copy Environment	-V	--export=[ALL   NONE   variables]		-V
Event Notification	-m abe	--mail-type=[events]	-B or -N	-m abe
Email Address	-M [address]	--mail-user=[address]	-u [address]	-M [address]
Job Name	-N [name]	--job-name=[name]	-J [name]	-N [name]
Job Restart	-r [y/n]	--requeue OR --no-requeue (NOTE: configurable default)	-r	-r [yes/no]
Working Directory	N/A	--workdir=[dir_name]	(submission directory)	-wd [directory]
Resource Sharing	-l naccesspolicy=singlejob	--exclusive OR --shared	-x	-l exclusive
Memory Size	-l mem=[MB]	--mem=[mem][M G T] OR --mem-per-cpu=[mem][M G T]	-M [MB]	-l mem_free=[memory][K M G]
Account to charge	-W group_list=[account]	--account=[account]	-P [account]	-A [account]
Tasks Per Node	-l mppnppn [PEs_per_node]	--tasks-per-node=[count]		(Fixed allocation_rule in PE)
CPUs Per Task		--cpus-per-task=[count]		
Job Dependency	-d [job_id]	--depend=[state:job_id]	-w [done   exit   finish]	-hold_jid [job_id   job_name]
Job Project		--wckey=[name]	-P [name]	-P [name]
Job host preference		--odelist=[nodes] AND/OR --exclude=[nodes]	-m [nodes]	-q [queue]@[node] OR -q [queue]@[hostgroup]
Quality Of Service	-l qos=[name]	--qos=[name]	J "name[array_spec]"	-t [array_spec]
Job Arrays	-t [array_spec]	--array=[array_spec] (Slurm version 2.6+)		-l [resource]=[value]
Generic Resources	-l other=[resource_spec]	--gres=[resource_spec]	-R "rusage[license_spec]"	-l [license]=[count]
Licenses		--licenses=[license_spec]		
Begin Time	-A "YYYY-MM-DD HH:MM:SS"	--begin=YYYY-MM-DD[THH:MM[:SS]]	-b[[year:]][month:]d[ay:]hour:minute	-a [YYMMDDhhmm]

Figure 11: Rosetta Mappings of Scheduler Commands from SchedMD

### A.3 Phases

Brief summary of Speed evolution phases.

#### A.3.1 Phase 4

Phase 4 had 7 SuperMicro servers with 4x A100 80GB GPUs each added, dubbed as “SPEED2”. We also moved from Grid Engine to SLURM.

#### A.3.2 Phase 3

Phase 3 had 4 vidpro nodes added from Dr. Amer totalling 6x P6 and 6x V100 GPUs added.

### A.3.3 Phase 2

Phase 2 saw 6x NVIDIA Tesla P6 added and 8x more compute nodes. The P6s replaced 4x of FirePro S7150.

### A.3.4 Phase 1

Phase 1 of Speed was of the following configuration:

- Sixteen, 32-core nodes, each with 512 GB of memory and approximately 1 TB of volatile-scratch disk space.
- Five AMD FirePro S7150 GPUs, with 8 GB of memory (compatible with the Direct X, OpenGL, OpenCL, and Vulkan APIs).

## B Frequently Asked Questions

### B.1 Where do I learn about Linux?

All Speed users are expected to have a basic understanding of Linux and its commonly used commands.

#### Software Carpentry

Software Carpentry provides free resources to learn software, including a workshop on the Unix shell. <https://software-carpentry.org/lessons/>

#### Udemy

There are a number of Udemy courses, including free ones, that will assist you in learning Linux. Active Concordia faculty, staff and students have access to Udemy courses. The course **Linux Mastery: Master the Linux Command Line in 11.5 Hours** is a good starting point for beginners. Visit <https://www.concordia.ca/it/services/udemy.html> to learn how Concordians may access Udemy.

### B.2 How to use the “bash shell” on Speed?

This section describes how to use the “bash shell” on Speed. Review Section 2.1.2 to ensure that your bash environment is set up.

#### B.2.1 How do I set bash as my login shell?

In order to set your default login shell to bash on Speed, your login shell on all GCS servers must be changed to bash. To make this change, create a ticket with the Service Desk (or email [help@concordia.ca](mailto:help@concordia.ca)) to request that bash become your default login shell for your ENCS user account on all GCS servers.

### B.2.2 How do I move into a bash shell on Speed?

To move to the bash shell, type **bash** at the command prompt. For example:

```
[speed-submit] [/home/a/a_user] > bash
bash-4.4$ echo $0
bash
```

**Note** how the command prompt changed from `[speed-submit] [/home/a/a_user] >` to `bash-4.4$` after entering the bash shell.

### B.2.3 How do I use the bash shell in an interactive session on Speed?

Below are examples of how to use **bash** as a shell in your interactive job sessions with both the **salloc** and **srun** commands.

- `salloc -ppt --mem=100G -N 1 -n 10 /encs/bin/bash`
- `srun --mem=50G -n 5 --pty /encs/bin/bash`

**Note:** Make sure the interactive job requests memory, cores, etc.

### B.2.4 How do I run scripts written in bash on Speed?

To execute bash scripts on Speed:

1. Ensure that the shebang of your bash job script is `#!/encs/bin/bash`
2. Use the **sbatch** command to submit your job script to the scheduler.

The Speed GitHub contains a sample bash job script.

## B.3 How to resolve “Disk quota exceeded” errors?

### B.3.1 Probable Cause

The “Disk quota exceeded” Error occurs when your application has run out of disk space to write to. On Speed this error can be returned when:

1. Your NFS-provided home is full and cannot be written to. You can verify this using **quota** and **bigfiles** commands.
2. The `/tmp` directory on the speed node your application is running on is full and cannot be written to.

### B.3.2 Possible Solutions

1. Use the `--chdir` job script option to set the directory that the job script is submitted from the **job working directory**. The **job working directory** is the directory that the job will write output files in.
2. The use local disk space is generally recommended for IO intensive operations. However, as the size of `/tmp` on speed nodes is 1TB it can be necessary for scripts to store temporary data elsewhere. Review the documentation for each module called within your script to determine how to set working directories for that application. The basic steps for this solution are:

- Review the documentation on how to set working directories for each module called by the job script.
- Create a working directory in speed-scratch for output files. For example, this command will create a subdirectory called **output** in your **speed-scratch** directory:

```
mkdir -m 750 /speed-scratch/$USER/output
```

- To create a subdirectory for recovery files:
- ```
mkdir -m 750 /speed-scratch/$USER/recovery
```
- Update the job script to write output to the subdirectories you created in your **speed-scratch** directory, e.g., `/speed-scratch/$USER/output`.

In the above example, `$USER` is an environment variable containing your ENCS username.

### B.3.3 Example of setting working directories for COMSOL

- Create directories for recovery, temporary, and configuration files. For example, to create these directories for your GCS ENCS user account:

```
mkdir -m 750 -p /speed-scratch/$USER/comsol/{recovery,tmp,config}
```

- Add the following command switches to the COMSOL command to use the directories created above:

```
-recoverydir /speed-scratch/$USER/comsol/recovery
-tmpdir /speed-scratch/$USER/comsol/tmp
-configuration/speed-scratch/$USER/comsol/config
```

In the above example, `$USER` is an environment variable containing your ENCS username.

### B.3.4 Example of setting working directories for Python Modules

By default when adding a python module the `/tmp` directory is set as the temporary repository for files downloads. The size of the `/tmp` directory on **speed-submit** is too small for pytorch. To add a python module

- Create your own tmp directory in your **speed-scratch** directory

```
mkdir /speed-scratch/$USER/tmp
```

- Use the tmp directory you created

```
setenv TMPDIR /speed-scratch/$USER/tmp
```

- Attempt the installation of pytorch

In the above example, `$USER` is an environment variable containing your ENCS username.

## B.4 How do I check my job's status?

When a job with a job id of 1234 is running or terminated, the status of that job can be tracked using 'sacct -j 1234'. squeue -j 1234 can show while the job is sitting in the queue as well. Long term statistics on the job after its terminated can be found using sstat -j 1234 after slurmctld purges its tracking state into the database.

## B.5 Why is my job pending when nodes are empty?

### B.5.1 Disabled nodes

It is possible that one or a number of the Speed nodes are disabled. Nodes are disabled if they require maintenance. To verify if Speed nodes are disabled, see if they are in a draining or drained state:

```
[serguei@speed-submit src] % sinfo --long --Node
Thu Oct 19 21:25:12 2023

```

| MODELIST | NODES | PARTITION | STATE   | CPUS | S:C:T  | MEMORY | TMP_DISK | WEIGHT | AVAIL_FE | REASON |
|----------|-------|-----------|---------|------|--------|--------|----------|--------|----------|--------|
| speed-01 | 1     | pa        | idle    | 32   | 2:16:1 | 257458 | 0        | 1      | gpu16    | none   |
| speed-03 | 1     | pa        | idle    | 32   | 2:16:1 | 257458 | 0        | 1      | gpu32    | none   |
| speed-05 | 1     | pg        | idle    | 32   | 2:16:1 | 515490 | 0        | 1      | gpu16    | none   |
| speed-07 | 1     | ps*       | mixed   | 32   | 2:16:1 | 515490 | 0        | 1      | cpu32    | none   |
| speed-08 | 1     | ps*       | drained | 32   | 2:16:1 | 515490 | 0        | 1      | cpu32    | UGE    |
| speed-09 | 1     | ps*       | drained | 32   | 2:16:1 | 515490 | 0        | 1      | cpu32    | UGE    |
| speed-10 | 1     | ps*       | drained | 32   | 2:16:1 | 515490 | 0        | 1      | cpu32    | UGE    |
| speed-11 | 1     | ps*       | idle    | 32   | 2:16:1 | 515490 | 0        | 1      | cpu32    | none   |
| speed-12 | 1     | ps*       | drained | 32   | 2:16:1 | 515490 | 0        | 1      | cpu32    | UGE    |
| speed-15 | 1     | ps*       | drained | 32   | 2:16:1 | 515490 | 0        | 1      | cpu32    | UGE    |
| speed-16 | 1     | ps*       | drained | 32   | 2:16:1 | 515490 | 0        | 1      | cpu32    | UGE    |
| speed-17 | 1     | pg        | drained | 32   | 2:16:1 | 515490 | 0        | 1      | gpu16    | UGE    |
| speed-19 | 1     | ps*       | idle    | 32   | 2:16:1 | 515490 | 0        | 1      | cpu32    | none   |
| speed-20 | 1     | ps*       | drained | 32   | 2:16:1 | 515490 | 0        | 1      | cpu32    | UGE    |
| speed-21 | 1     | ps*       | drained | 32   | 2:16:1 | 515490 | 0        | 1      | cpu32    | UGE    |
| speed-22 | 1     | ps*       | drained | 32   | 2:16:1 | 515490 | 0        | 1      | cpu32    | UGE    |
| speed-23 | 1     | ps*       | idle    | 32   | 2:16:1 | 515490 | 0        | 1      | cpu32    | none   |
| speed-24 | 1     | ps*       | idle    | 32   | 2:16:1 | 515490 | 0        | 1      | cpu32    | none   |
| speed-25 | 1     | pg        | idle    | 32   | 2:16:1 | 257458 | 0        | 1      | gpu32    | none   |
| speed-25 | 1     | pa        | idle    | 32   | 2:16:1 | 257458 | 0        | 1      | gpu32    | none   |
| speed-27 | 1     | pg        | idle    | 32   | 2:16:1 | 257458 | 0        | 1      | gpu32    | none   |
| speed-27 | 1     | pa        | idle    | 32   | 2:16:1 | 257458 | 0        | 1      | gpu32    | none   |
| speed-29 | 1     | ps*       | idle    | 32   | 2:16:1 | 515490 | 0        | 1      | cpu32    | none   |
| speed-30 | 1     | ps*       | drained | 32   | 2:16:1 | 515490 | 0        | 1      | cpu32    | UGE    |
| speed-31 | 1     | ps*       | drained | 32   | 2:16:1 | 515490 | 0        | 1      | cpu32    | UGE    |
| speed-32 | 1     | ps*       | drained | 32   | 2:16:1 | 515490 | 0        | 1      | cpu32    | UGE    |
| speed-33 | 1     | ps*       | idle    | 32   | 2:16:1 | 515490 | 0        | 1      | cpu32    | none   |
| speed-34 | 1     | ps*       | idle    | 32   | 2:16:1 | 515490 | 0        | 1      | cpu32    | none   |
| speed-35 | 1     | ps*       | drained | 32   | 2:16:1 | 515490 | 0        | 1      | cpu32    | UGE    |
| speed-36 | 1     | ps*       | drained | 32   | 2:16:1 | 515490 | 0        | 1      | cpu32    | UGE    |
| speed-37 | 1     | pt        | idle    | 256  | 2:64:2 | 980275 | 0        | 1      | gpu20,mi | none   |
| speed-38 | 1     | pt        | idle    | 256  | 2:64:2 | 980275 | 0        | 1      | gpu20,mi | none   |
| speed-39 | 1     | pt        | idle    | 256  | 2:64:2 | 980275 | 0        | 1      | gpu20,mi | none   |
| speed-40 | 1     | pt        | idle    | 256  | 2:64:2 | 980275 | 0        | 1      | gpu20,mi | none   |
| speed-41 | 1     | pt        | idle    | 256  | 2:64:2 | 980275 | 0        | 1      | gpu20,mi | none   |
| speed-42 | 1     | pt        | idle    | 256  | 2:64:2 | 980275 | 0        | 1      | gpu20,mi | none   |
| speed-43 | 1     | pt        | idle    | 256  | 2:64:2 | 980275 | 0        | 1      | gpu20,mi | none   |

Note which nodes are in the state of **drained**. Why the state is drained can be found in the reason column.

Your job will run once an occupied node becomes available or the maintenance has been completed and the disabled nodes have a state of **idle**.

### B.5.2 Error in job submit request.

It is possible that your job is pending, because the job requested resources that are not available within Speed. To verify why job id 1234 is not running, execute ‘`sacct -j 1234`’. A summary of the reasons is available via the `squeue` command.

## C Sister Facilities

Below is a list of resources and facilities similar to Speed at various capacities. Depending on your research group and needs, they might be available to you. They are not managed by HPC/NAG of AITS, so contact their respective representatives.

- `computation.encs` CPU only 3-machine cluster running longer jobs without a scheduler at the moment
- `apini.encs` cluster for teaching and MPI programming (see the corresponding course in CSSE)
- Computer Science and Software Engineering (CSSE) Virya GPU Cluster. For CSSE members only. The cluster has 4 nodes with total of 32 NVIDIA GPUs (a mix of V100s and A100s). To request access send email to `virya.help AT concordia.ca`.
- Dr. Maria Amer’s VidPro group’s nodes in Speed (-01, -03, -25, -27) with additional V100 and P6 GPUs.
- There are various Lambda Labs other GPU servers and like computers acquired by individual researchers; if you are member of their research group, contact them directly. These resources are not managed by us.
  - Dr. Amin Hammad’s `construction.encs` Lambda Labs station
  - Dr. Hassan Rivaz’s `impactlab.encs` Lambda Labs station
  - Dr. Nizar Bouguila’s `xailab.encs` Lambda Labs station
  - Dr. Roch Glitho’s `femto.encs` server
  - Dr. Maria Amer’s `venom.encs` Lambda Labs station
  - Dr. Leon Wang’s `guerrera.encs` DGX station
- Dr. Ivan Contreras’ servers (managed by AITS)
- If you are a member of School of Health (formerly PERFORM Center), you may have access to their local PERFORM’s High Performance Computing (HPC) Cluster. Contact Thomas Beaudry for details and how to obtain access.
- All Concordia students have access to the Library’s small Technology Sandbox testing cluster that also runs Slurm. Email `sean.cooney AT concordia.ca` for details.
- Digital Research Alliance Canada (Compute Canada / Calcul Quebec), <https://alliancecan.ca/>. Follow this link on the information how to obtain access (students need to be sponsored by their supervising faculty members, who should create accounts first). Their SLURM examples are here: [https://docs.alliancecan.ca/wiki/Running\\_jobs](https://docs.alliancecan.ca/wiki/Running_jobs)

## References

- [1] 3DS. Abaqus. [online], 2019–2021. <https://www.3ds.com/products-services/simulia/products/abaqus/>.
- [2] ANSYS. FLUENT. [online], 2000–2012. <http://www.ansys.com/Products/Simulation+Technology/Fluid+Dynamics/ANSYS+FLUENT>.
- [3] Belkacem Belabes and Marius Paraschivoiu. Numerical study of the effect of turbulence intensity on VAWT performance. *Energy*, 233:121139, 2021. <https://doi.org/10.1016/j.energy.2021.121139>.
- [4] Belkacem Belabes and Marius Paraschivoiu. CFD study of the aerodynamic performance of a vertical axis wind turbine in the wake of another turbine. In *Proceedings of the CSME International Congress*, 2022. <https://doi.org/10.7939/r3-rker-1746>.
- [5] Belkacem Belabes and Marius Paraschivoiu. CFD modeling of vertical-axis wind turbine wake interaction. *Transactions of the Canadian Society for Mechanical Engineering*, pages 1–10, 2023. <https://doi.org/10.1139/tcsme-2022-0149>.
- [6] Amine Boukhtouta, Nour-Eddine Lakhdari, Serguei A. Mokhov, and Mourad Debbabi. Towards fingerprinting malicious traffic. In *Proceedings of ANT’13*, volume 19, pages 548–555. Elsevier, June 2013.
- [7] Amy Brown and Greg Wilson, editors. *The Architecture of Open Source Applications: Elegance, Evolution, and a Few Fearless Hacks*, volume I. aosabook.org, March 2012. Online at <http://aosabook.org>.
- [8] Goutam Yelluru Gopal and Maria Amer. Mobile vision transformer-based visual object tracking. In *34th British Machine Vision Conference (BMVC)*, Aberdeen, UK, November 2023. <https://arxiv.org/abs/2309.05829> and <https://github.com/goutamyg/MVT>.
- [9] Goutam Yelluru Gopal and Maria Amer. Separable self and mixed attention transformers for efficient object tracking. In *IEEE/CVF Winter Conference on Applications of Computer Vision (WACV)*, Waikoloa, Hawaii, January 2024. <https://arxiv.org/abs/2309.03979> and <https://github.com/goutamyg/SMAT>.
- [10] Haotao Lai. An OpenISS framework specialization for deep learning-based person re-identification. Master’s thesis, Department of Computer Science and Software Engineering, Concordia University, Montreal, Canada, August 2019. <https://spectrum.library.concordia.ca/id/eprint/985788/>.
- [11] Haotao Lai et al. OpenISS keras-yolo3 v0.1.0, June 2021. <https://github.com/OpenISS/openiss-yolov3>.
- [12] Haotao Lai et al. Openiss person re-identification baseline v0.1.1, June 2021. <https://github.com/OpenISS/openiss-reid-tfk>.
- [13] MathWorks. MATLAB. [online], 2000–2012. <http://www.mathworks.com/products/matlab/>.
- [14] Serguei Mokhov, Jonathan Llewellyn, Carlos Alarcon Meza, Tariq Daradkeh, and Gillian Roper. The use of containers in OpenGL, ML and HPC for teaching and research support. In *ACM SIGGRAPH 2023 Posters*, SIGGRAPH ’23, New York, NY, USA, 2023. ACM. <https://doi.org/10.1145/3588028.3603676>.
- [15] Serguei A. Mokhov. The use of machine learning with signal- and NLP processing of source code to fingerprint, detect, and classify vulnerabilities and weaknesses with MARFCAT. Technical Report NIST SP 500-283, NIST, October 2011. Report: [http://www.nist.gov/manuscript-publication-search.cfm?pub\\_id=909407](http://www.nist.gov/manuscript-publication-search.cfm?pub_id=909407), online e-print at <http://arxiv.org/abs/1010.2511>.
- [16] Serguei A. Mokhov. *Intensional Cyberforensics*. PhD thesis, Department of Computer Science and Software Engineering, Concordia University, Montreal, Canada, September 2013. Online at <http://arxiv.org/abs/1312.0466>.
- [17] Serguei A. Mokhov, Michael J. Assels, Joey Paquet, and Mourad Debbabi. Automating MAC



- spoofers evidence gathering and encoding for investigations. In Frederic Cuppens et al., editors, *Proceedings of The 7th International Symposium on Foundations & Practice of Security (FPS'14)*, LNCS 8930, pages 168–183. Springer, November 2014. Full paper.
- [18] Serguei A. Mokhov, Michael J. Assels, Joey Paquet, and Mourad Debbabi. Toward automated MAC spoofers investigations. In *Proceedings of C3S2E'14*, pages 179–184. ACM, August 2014. Short paper.
- [19] Serguei A. Mokhov and Scott Bunnell. Speed server farm: Gina Cody School of ENCS HPC facility. [online], 2018–2019. [https://docs.google.com/presentation/d/1bWbGQvYsuJ4U2WsfLYp8S3yb4i70dU7QDn3l\\_Q9mYis](https://docs.google.com/presentation/d/1bWbGQvYsuJ4U2WsfLYp8S3yb4i70dU7QDn3l_Q9mYis).
- [20] Serguei A. Mokhov, Joey Paquet, and Mourad Debbabi. The use of NLP techniques in static code analysis to detect weaknesses and vulnerabilities. In Maria Sokolova and Peter van Beek, editors, *Proceedings of Canadian Conference on AI'14*, volume 8436 of *LNAI*, pages 326–332. Springer, May 2014. Short paper.
- [21] Parna Niksirat, Adriana Dacu, and Krzysztof Skonieczny. The effects of reduced-gravity on planetary rover mobility. *International Journal of Robotics Research*, 39(7):797–811, 2020. <https://doi.org/10.1177/0278364920913945>.
- [22] Chet Ramey. The Bourne-Again Shell. In Brown and Wilson [7]. <http://aosabook.org/en/bash.html>.
- [23] Rob Schreiber. MATLAB. *Scholarpedia*, 2(6):2929, 2007. <http://www.scholarpedia.org/article/MATLAB>.
- [24] The MARF Research and Development Group. The Modular Audio Recognition Framework and its Applications. [online], 2002–2014. <http://marf.sf.net> and <http://arxiv.org/abs/0905.1235>, last viewed May 2015.